

Step 1: Import Required Libraries

This cell imports all necessary libraries for data manipulation, visualization, and modeling.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
```

Step 2: Load the Data

Load all sheets from the Excel file and display a preview of key datasets.

```
# Step 2: Load the Data

# Define file path
file_path = 'HrData.xlsx'

# Load sheets into DataFrames
employee_df = pd.read_excel(file_path, sheet_name='Employee')
performance_df = pd.read_excel(file_path,
                                sheet_name='PerformanceRating')
education_level_df = pd.read_excel(file_path,
                                    sheet_name='EducationLevel')
rating_level_df = pd.read_excel(file_path, sheet_name='RatingLevel')
satisfied_level_df = pd.read_excel(file_path,
                                    sheet_name='SatisfiedLevel')

# Convert ReviewDate to datetime
performance_df['ReviewDate'] =
pd.to_datetime(performance_df['ReviewDate'])

# Display previews
print('Employee Data:')
print(employee_df.head())

print('\nPerformance Rating Data:')
print(performance_df.head())
```

Employee Data:

	EmployeeID	FirstName	LastName	Gender	Age	BusinessTravel	\
0	3012-1A41	Leonelle	Simco	Female	30	Some Travel	
1	CBCB-9C9D	Leonerd	Aland	Male	38	Some Travel	
2	95D7-1CE9	Ahmed	Sykes	Male	43	Some Travel	

3	47A0-559B	Ermentrude	Berrie	Non-Binary	39	Some Travel
4	42CC-040A	Stace	Savege	Female	29	Some Travel

	Department	DistanceFromHome	State
Ethnicity ... \			
0	Sales	27	IL
White ...			
1	Sales	23	CA
White ...			
2	Human Resources	29	CA Asian or Asian
American ...			
3	Technology	12	IL
White ...			
4	Human Resources	29	CA
White ...			

	MaritalStatus	Salary	StockOptionLevel	OverTime	HireDate
Attrition \					
0	Divorced	102059	1	No	2012-01-03
No					
1	Single	157718	0	Yes	2012-01-04
No					
2	Married	309964	1	No	2012-01-04
No					
3	Married	293132	0	No	2012-01-05
No					
4	Single	49606	0	No	2012-01-05
Yes					

	YearsAtCompany	YearsInMostRecentRole	YearsSinceLastPromotion	\
0	10	4	9	
1	10	6	10	
2	10	6	10	
3	10	10	10	
4	6	1	1	

	YearsWithCurrManager
0	7
1	0
2	8
3	0
4	6

[5 rows x 23 columns]

Performance Rating Data:

	PerformanceID	EmployeeID	ReviewDate	EnvironmentSatisfaction	\
0	PR01	79F7-78EC	2013-02-01	5	
1	PR02	B61E-0F26	2013-03-01	5	
2	PR03	F5E3-48BB	2013-03-01	3	

3	PR04	0678-748A	2013-04-01	5
4	PR05	541F-3E19	2013-04-01	5

	JobSatisfaction	RelationshipSatisfaction
TrainingOpportunitiesWithinYear \		

0	4	5
1		
1	4	4
1		
2	4	5
3		
3	3	2
2		
4	2	3
1		

	TrainingOpportunitiesTaken	WorkLifeBalance	SelfRating
ManagerRating \			

0	0	4	4
4			
1	1	4	4
3			
2	2	3	5
4			
3	0	2	3
2			
4	0	4	4
3			

	DataQualityStatus
--	-------------------

0	Valid
1	Valid
2	Valid
3	Valid
4	Valid

Step 3: Performance GAP Analysis

Objective: Identify the gap between current and desired performance. Here we use the latest performance review per employee and compare each employee's ManagerRating to the average for their JobRole.

```
# Get the most recent performance review per employee
latest_performance =
performance_df.loc[performance_df.groupby('EmployeeID')
['ReviewDate'].idxmax()]
```

```

# Merge the latest performance review with employee data
merged_df = pd.merge(employee_df, latest_performance, on='EmployeeID',
how='left')

# Calculate average ManagerRating per JobRole
avg_rating_by_role = merged_df.groupby('JobRole')
['ManagerRating'].mean()

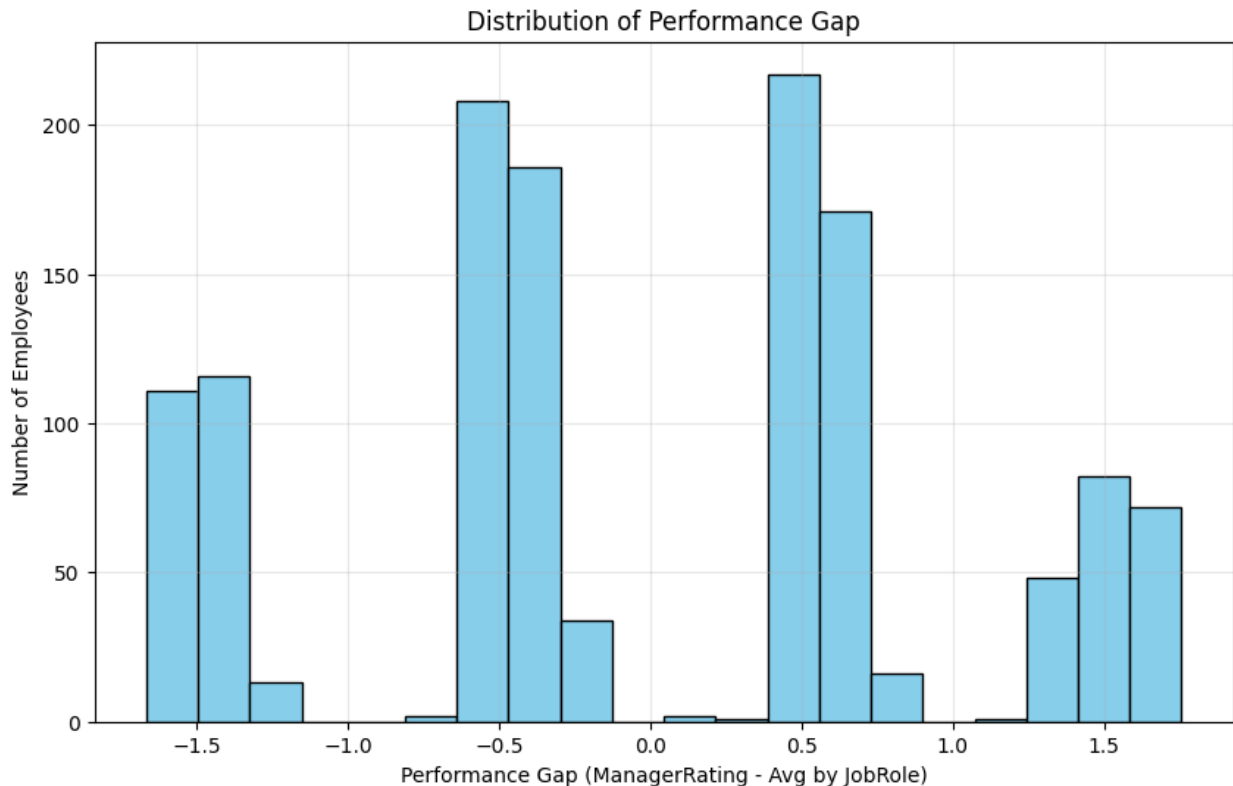
# Compute the performance gap per employee (ManagerRating - average
for their job role)
merged_df['Gap'] = merged_df['ManagerRating'] -
merged_df['JobRole'].map(avg_rating_by_role)

# Visualize the distribution of performance gaps
plt.figure(figsize=(10, 6))
plt.hist(merged_df['Gap'].dropna(), bins=20, color='skyblue',
edgecolor='black')
plt.title('Distribution of Performance Gap')
plt.xlabel('Performance Gap (ManagerRating - Avg by JobRole)')
plt.ylabel('Number of Employees')
plt.grid(True, alpha=0.3)
plt.show()

# Calculate percentage of employees with a negative gap
negative_gap_pct = (merged_df['Gap'] < 0).mean() * 100
print(f"Percentage of employees performing below their job role
average: {negative_gap_pct:.2f}%")

# Explore correlation with training opportunities taken (if column
exists in performance_df)
if 'TrainingOpportunitiesTaken' in merged_df.columns:
    correlation =
merged_df['Gap'].corr(merged_df['TrainingOpportunitiesTaken'])
    print(f"Correlation between Performance Gap and
TrainingOpportunitiesTaken: {correlation:.2f}")
else:
    print("Column 'TrainingOpportunitiesTaken' not found in merged
data.")

```



Percentage of employees performing below their job role average:
45.58%

Correlation between Performance Gap and TrainingOpportunitiesTaken:
0.01

Step 4: Training GAP Analysis

Objective: Calculate the gap between available and attended training opportunities and identify the top 10 employees with the largest training gaps.

Step 4: Training GAP Analysis

*# Objective: Calculate the gap between available and attended training opportunities
and identify the top 10 employees with the largest training gaps.*

```
# Calculate the training gap
performance_df['TrainingGap'] =
performance_df['TrainingOpportunitiesWithinYear'] -
performance_df['TrainingOpportunitiesTaken']
```

```
# Summarize training gaps
gap_summary = performance_df['TrainingGap'].describe()
```

```

print('\nTraining Gap Statistics:\n', gap_summary)

# Extract employees with the largest training gaps (Top 10)
top_training_gaps = performance_df.sort_values(by='TrainingGap',
ascending=False).head(10)

# Retrieve their details from the Employee table
df_employee_selected =
employee_df[employee_df['EmployeeID'].isin(top_training_gaps['Employee
ID'])]

# Merge performance data with employee details
df_gap_analysis = pd.merge(top_training_gaps, df_employee_selected,
on='EmployeeID', how='left')

# Select important columns
df_gap_analysis = df_gap_analysis[['EmployeeID', 'FirstName',
'LastName', 'Department',
'TrainingOpportunitiesWithinYear',
'TrainingOpportunitiesTaken', 'TrainingGap']]

# Display the top 10 employees in need of training
print('\nTop 10 Employees in Need of Training:\n', df_gap_analysis)

```

Training Gap Statistics:

```

count      6709.000000
mean        1.147116
std         1.014636
min         0.000000
25%         0.000000
50%         1.000000
75%         2.000000
max         3.000000

```

Name: TrainingGap, dtype: float64

Top 10 Employees in Need of Training:

	EmployeeID	FirstName	LastName	Department	\
0	3EB2-9111	Ginger	Blinde	Technology	
1	2E72-4BF1	Grace	Gohier	Sales	
2	5468-EEE1	Forbes	Toretta	Technology	
3	10E9-4C86	Ignacius	Dockrill	Technology	
4	D676-4ECC	Florenza	Nesbit	Technology	
5	2B7A-9C73	Bryanty	Wickersley	Technology	
6	B6EC-313E	Corbin	Goody	Human Resources	
7	7749-B277	Caryl	Roycroft	Human Resources	
8	1799-5B3F	Howey	Woolis	Sales	
9	819A-2C9C	Hillary	Atchly	Technology	

TrainingOpportunitiesWithinYear TrainingOpportunitiesTaken

TrainingGap		
0	3	0
3		
1	3	0
3		
2	3	0
3		
3	3	0
3		
4	3	0
3		
5	3	0
3		
6	3	0
3		
7	3	0
3		
8	3	0
3		
9	3	0
3		

Step 5: 9-Box Grid Analysis

Objective: Place employees in a 3x3 grid based on their current performance and future potential. For current performance we use ManagerRating; for future potential, we compute the trend (slope) of ManagerRating over time using linear regression.

```
# Function to calculate the slope of ManagerRating over time for each employee
def calculate_slope(group):
    if len(group) < 2: # Need at least 2 points to calculate a trend
        return np.nan
    X = (group['ReviewDate'] -
group['ReviewDate'].min()).dt.days.values.reshape(-1, 1)
    y = group['ManagerRating'].values
    model = LinearRegression().fit(X, y)
    return model.coef_[0]

# Calculate slopes for each employee based on their performance reviews
slopes = performance_df.groupby('EmployeeID').apply(calculate_slope)
merged_df['Slope'] = merged_df['EmployeeID'].map(slopes)

# Categorize current performance based on ManagerRating
def categorize_performance(rating):
    if pd.isna(rating):
```

```

        return np.nan
    elif rating <= 2:
        return 'Low'
    elif rating == 3:
        return 'Average'
    else:
        return 'High'

merged_df['PerformanceCategory'] =
merged_df['ManagerRating'].apply(categorize_performance)

# Filter employees with valid slope values and then categorize
potential into 3 quantiles
valid_df = merged_df.dropna(subset=['Slope'])
valid_df['PotentialCategory'] = pd.qcut(valid_df['Slope'], 3,
labels=['Low', 'Average', 'High'])

# Create the 9-box grid as a crosstab of Performance vs. Potential
grid = pd.crosstab(valid_df['PerformanceCategory'],
valid_df['PotentialCategory'])
print('9-Box Grid Distribution:')
print(grid)

# Visualize the 9-box grid using a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(grid, annot=True, fmt='d', cmap='Blues')
plt.title('9-Box Grid: Performance vs. Potential')
plt.xlabel('Future Potential')
plt.ylabel('Current Performance')
plt.show()

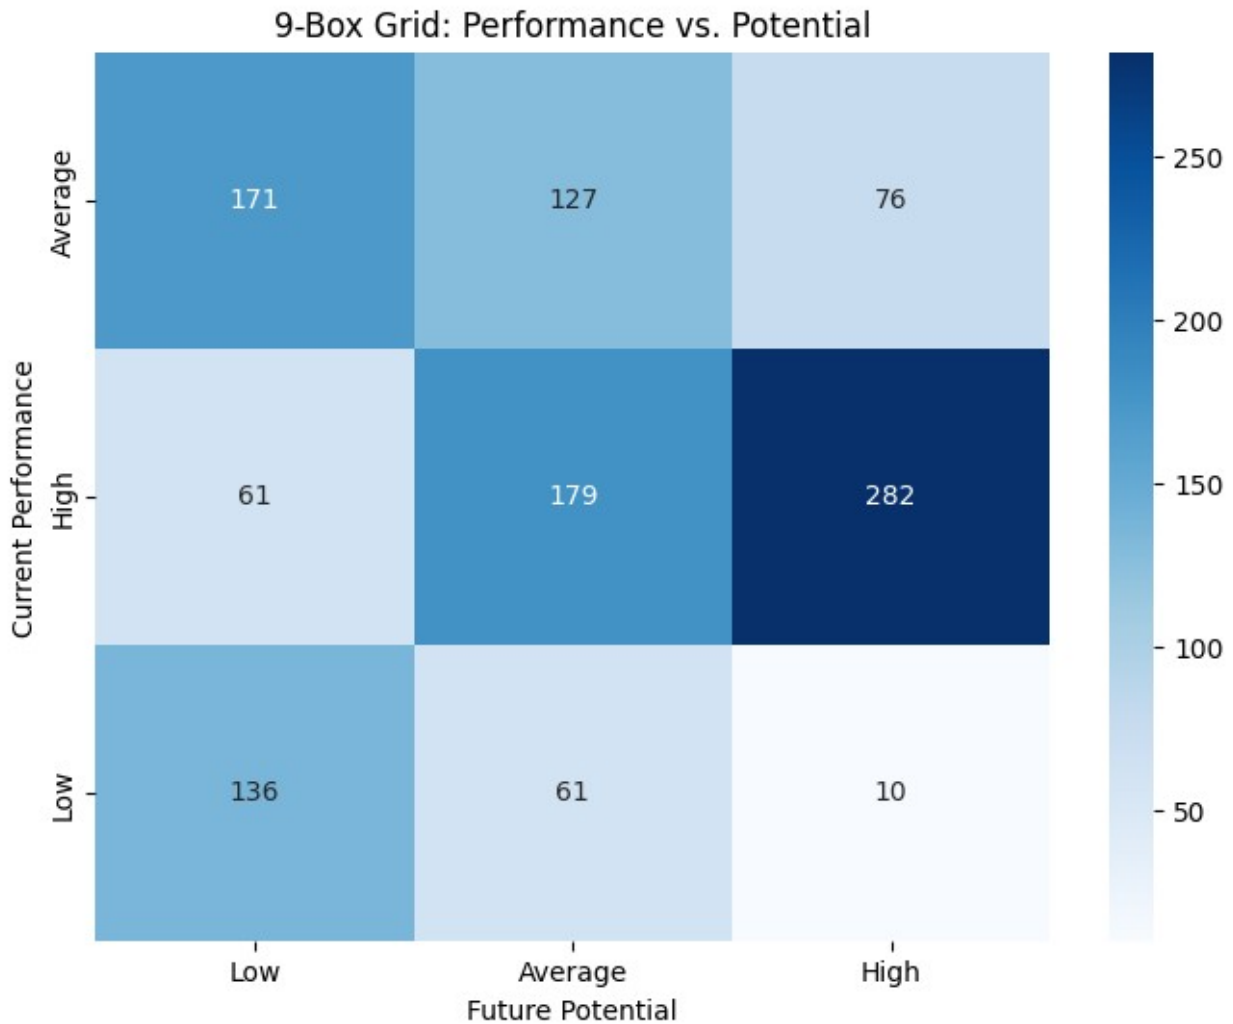
/tmp/ipykernel_872422/3655308779.py:11: DeprecationWarning:
DataFrameGroupBy.apply operated on the grouping columns. This behavior
is deprecated, and in a future version of pandas the grouping columns
will be excluded from the operation. Either pass
`include_groups=False` to exclude the groupings or explicitly select
the grouping columns after groupby to silence this warning.
    slopes = performance_df.groupby('EmployeeID').apply(calculate_slope)
/tmp/ipykernel_872422/3655308779.py:29: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
returning-a-view-versus-a-copy
    valid_df['PotentialCategory'] = pd.qcut(valid_df['Slope'], 3,
labels=['Low', 'Average', 'High'])

9-Box Grid Distribution:
PotentialCategory    Low  Average  High

```


PerformanceCategory			
Average	171	127	76
High	61	179	282
Low	136	61	10



Step 6: KPIs & Turnover Analysis

Objective: Calculate key HR metrics including employee turnover rate, average satisfaction, and average tenure. Then, analyze turnover by department, salary, and experience.

```
# 1. Employee Turnover Rate
turnover_rate =
employee_df['Attrition'].value_counts(normalize=True).get('Yes', 0) *
100
```

```
# 2. Convert satisfaction levels to numerical values
```

```

satisfaction_mapping = {
    "Very Dissatisfied": 1,
    "Dissatisfied": 2,
    "Neutral": 3,
    "Satisfied": 4,
    "Very Satisfied": 5
}
satisfied_level_df['SatisfactionLevelNumeric'] =
satisfied_level_df['SatisfactionLevel'].map(satisfaction_mapping)

# Calculate the average employee satisfaction
avg_satisfaction =
satisfied_level_df['SatisfactionLevelNumeric'].mean()

# 3. Average tenure of employees in the company
avg_hiring_time = employee_df['YearsAtCompany'].mean()

# Aggregate KPI Results
kpi_results = {
    "Turnover Rate (%)": turnover_rate,
    "Average Satisfaction Level": avg_satisfaction,
    "Average Hiring Time (Years)": avg_hiring_time
}

# Turnover Analysis
df_turnover = employee_df[employee_df['Attrition'] == 'Yes']
turnover_by_department = df_turnover['Department'].value_counts()
turnover_by_income = df_turnover['Salary'].median()
turnover_by_experience = df_turnover['YearsAtCompany'].median()

turnover_results = {
    "Turnover by Department": turnover_by_department.to_dict(),
    "Median Salary of Ex-Employees": turnover_by_income,
    "Median Experience of Ex-Employees (Years)":
turnover_by_experience
}

# Print KPIs Analysis results
print('=== KPIs Analysis ===')
for key, value in kpi_results.items():
    print(f"{key}: {value}")

print('\n=== Turnover Analysis ===')
for key, value in turnover_results.items():
    print(f"{key}: {value}")

=== KPIs Analysis ===
Turnover Rate (%): 16.122448979591837
Average Satisfaction Level: 2.75
Average Hiring Time (Years): 4.562585034013606

```

```
=== Turnover Analysis ===  
Turnover by Department: {'Technology': 133, 'Sales': 92, 'Human  
Resources': 12}  
Median Salary of Ex-Employees: 50660.0  
Median Experience of Ex-Employees (Years): 1.0
```

Step 7: Balanced Scorecard Analysis

Objective: Evaluate HR performance across multiple perspectives:

- **Financial Performance:** Attrition rate as a cost indicator.
- **Employee Satisfaction:** Average satisfaction scores.
- **Internal Processes:** Training utilization and promotion frequency.
- **Learning and Growth:** Training participation and recent promotions.

```
# Financial Perspective: Attrition Rate  
attrition_rate = (employee_df['Attrition'] == 'Yes').mean() * 100  
print(f"Attrition Rate: {attrition_rate:.2f}%")  
  
# Employee Satisfaction: Average satisfaction scores  
satisfaction_cols = ['EnvironmentSatisfaction', 'JobSatisfaction',  
                    'RelationshipSatisfaction', 'WorkLifeBalance']  
  
# Some satisfaction columns might be in the performance_df or  
# employee_df; adjust as needed  
if set(satisfaction_cols).issubset(latest_performance.columns):  
    latest_performance['OverallSatisfaction'] =  
        latest_performance[satisfaction_cols].mean(axis=1)  
    overall_satisfaction =  
        latest_performance['OverallSatisfaction'].mean()  
    print(f"Overall Average Satisfaction (1-5 scale):  
{overall_satisfaction:.2f}")  
else:  
    print("Satisfaction columns not found in latest performance  
data.")  
  
# Internal Processes: Training Utilization and Years Since Last  
# Promotion  
if 'TrainingOpportunitiesWithinYear' in latest_performance.columns and  
    'TrainingOpportunitiesTaken' in latest_performance.columns:  
    latest_performance['TrainingUtilization'] =  
        (latest_performance['TrainingOpportunitiesTaken'] /  
         latest_performance['TrainingOpportunitiesWithinYear'])  
    avg_training_utilization =  
        latest_performance['TrainingUtilization'].mean()  
    print(f"Average Training Utilization:  
{avg_training_utilization:.2f}")
```

```

else:
    print("Training data not found in latest performance data.")

if 'YearsSinceLastPromotion' in employee_df.columns:
    avg_years_since_promo =
employee_df['YearsSinceLastPromotion'].mean()
    print(f"Average Years Since Last Promotion:
{avg_years_since_promo:.2f}")
else:
    print("Column 'YearsSinceLastPromotion' not found in employee
data.")

# Learning and Growth: Training Opportunities Taken and Recent
Promotions
if 'TrainingOpportunitiesTaken' in latest_performance.columns:
    avg_training_taken =
latest_performance['TrainingOpportunitiesTaken'].mean()
    print(f"Average Training Opportunities Taken:
{avg_training_taken:.2f}")
else:
    print("Column 'TrainingOpportunitiesTaken' not found in latest
performance data.")

if 'YearsSinceLastPromotion' in employee_df.columns:
    recent_promo_pct = (employee_df['YearsSinceLastPromotion'] <=
1).mean() * 100
    print(f"Percentage of Employees with Promotion in Last Year:
{recent_promo_pct:.2f}%")
else:
    print("Column 'YearsSinceLastPromotion' not found in employee
data.")

# Optional: Visualize key metrics
metrics = {
    'Attrition Rate (%)': attrition_rate,
    'Overall Satisfaction': overall_satisfaction if
'overall_satisfaction' in locals() else np.nan,
    'Training Utilization': avg_training_utilization if
'avg_training_utilization' in locals() else np.nan,
    'Avg Years Since Promo': avg_years_since_promo if
'avg_years_since_promo' in locals() else np.nan,
    'Avg Training Taken': avg_training_taken if 'avg_training_taken'
in locals() else np.nan,
    'Recent Promo (%)': recent_promo_pct if 'recent_promo_pct' in
locals() else np.nan
}

plt.figure(figsize=(12, 6))
plt.bar(metrics.keys(), metrics.values(), color='teal')
plt.title('Balanced Scorecard Metrics')

```

```
plt.ylabel('Value')
plt.xticks(rotation=45)
plt.show()
```

Attrition Rate: 16.12%

Overall Average Satisfaction (1-5 scale): 3.46

Average Training Utilization: 0.46

Average Years Since Last Promotion: 3.44

Average Training Opportunities Taken: 0.83

Percentage of Employees with Promotion in Last Year: 34.83%

