

```
# Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier,
RandomForestRegressor
from sklearn.metrics import accuracy_score, classification_report,
mean_squared_error
```

Step 1: Loading the Data

```
# Load data from Excel sheets
file_path = 'HrData.xlsx'
employee_df = pd.read_excel(file_path, sheet_name='Employee')
performance_df = pd.read_excel(file_path,
sheet_name='PerformanceRating')

# Convert date columns to datetime
employee_df['HireDate'] = pd.to_datetime(employee_df['HireDate'])
performance_df['ReviewDate'] =
pd.to_datetime(performance_df['ReviewDate'])

# Display sample data to verify loading
print("Employee Data Sample:")
print(employee_df.head())
print("\nPerformance Rating Data Sample:")
print(performance_df.head())
```

Employee Data Sample:

	EmployeeID	FirstName	LastName	Gender	Age	BusinessTravel	\
0	3012-1A41	Leonelle	Simco	Female	30	Some Travel	
1	CBCB-9C9D	Leonerd	Aland	Male	38	Some Travel	
2	95D7-1CE9	Ahmed	Sykes	Male	43	Some Travel	
3	47A0-559B	Ermentrude	Berrie	Non-Binary	39	Some Travel	
4	42CC-040A	Stace	Savege	Female	29	Some Travel	

	Department	DistanceFromHome	State
Ethnicity ... \			
0	Sales	27	IL
White ...			
1	Sales	23	CA
White ...			
2	Human Resources	29	CA Asian or Asian
American ...			
3	Technology	12	IL
White ...			
4	Human Resources	29	CA
White ...			

	MaritalStatus	Salary	StockOptionLevel	OverTime	HireDate
Attrition \					
0	Divorced	102059	1	No	2012-01-03
No					
1	Single	157718	0	Yes	2012-01-04
No					
2	Married	309964	1	No	2012-01-04
No					
3	Married	293132	0	No	2012-01-05
No					
4	Single	49606	0	No	2012-01-05
Yes					

	YearsAtCompany	YearsInMostRecentRole	YearsSinceLastPromotion	\
0	10	4	9	
1	10	6	10	
2	10	6	10	
3	10	10	10	
4	6	1	1	

	YearsWithCurrManager
0	7
1	0
2	8
3	0
4	6

[5 rows x 23 columns]

Performance Rating Data Sample:

	PerformanceID	EmployeeID	ReviewDate	EnvironmentSatisfaction	\
0	PR01	79F7-78EC	2013-02-01	5	
1	PR02	B61E-0F26	2013-03-01	5	
2	PR03	F5E3-48BB	2013-03-01	3	
3	PR04	0678-748A	2013-04-01	5	
4	PR05	541F-3E19	2013-04-01	5	

	JobSatisfaction	RelationshipSatisfaction
TrainingOpportunitiesWithinYear \		
0	4	5
1		
1	4	4
1		
2	4	5
3		
3	3	2
2		
4	2	3
1		

ManagerRating \	TrainingOpportunitiesTaken	WorkLifeBalance	SelfRating
0	0	4	4
4			
1	1	4	4
3			
2	2	3	5
4			
3	0	2	3
2			
4	0	4	4
3			

DataQualityStatus
0 Valid
1 Valid
2 Valid
3 Valid
4 Valid

Analysis 1: Employee Engagement Analysis

Objective: This analysis measures employee satisfaction and motivation using the EnvironmentSatisfaction column from performance data and correlates it with performance metrics (ManagerRating). It helps identify areas for improvement, such as the work environment or corporate culture, if satisfaction levels are low.

```
# Get the latest performance review per employee
latest_performance =
performance_df.loc[performance_df.groupby('EmployeeID')
['ReviewDate'].idxmax()]

# Merge with employee data
merged_df = pd.merge(employee_df, latest_performance, on='EmployeeID',
how='left')

# Calculate average environmental satisfaction
env_satisfaction_avg = merged_df['EnvironmentSatisfaction'].mean()
print(f"Average Environmental Satisfaction (1-5 scale):
{env_satisfaction_avg:.2f}")

# Visualize the distribution of Environmental Satisfaction
plt.figure(figsize=(10, 6))
plt.hist(merged_df['EnvironmentSatisfaction'].dropna(), bins=5,
color='skyblue', edgecolor='black')
plt.title('Distribution of Environmental Satisfaction')
plt.xlabel('Satisfaction Level (1-5)')
```

```

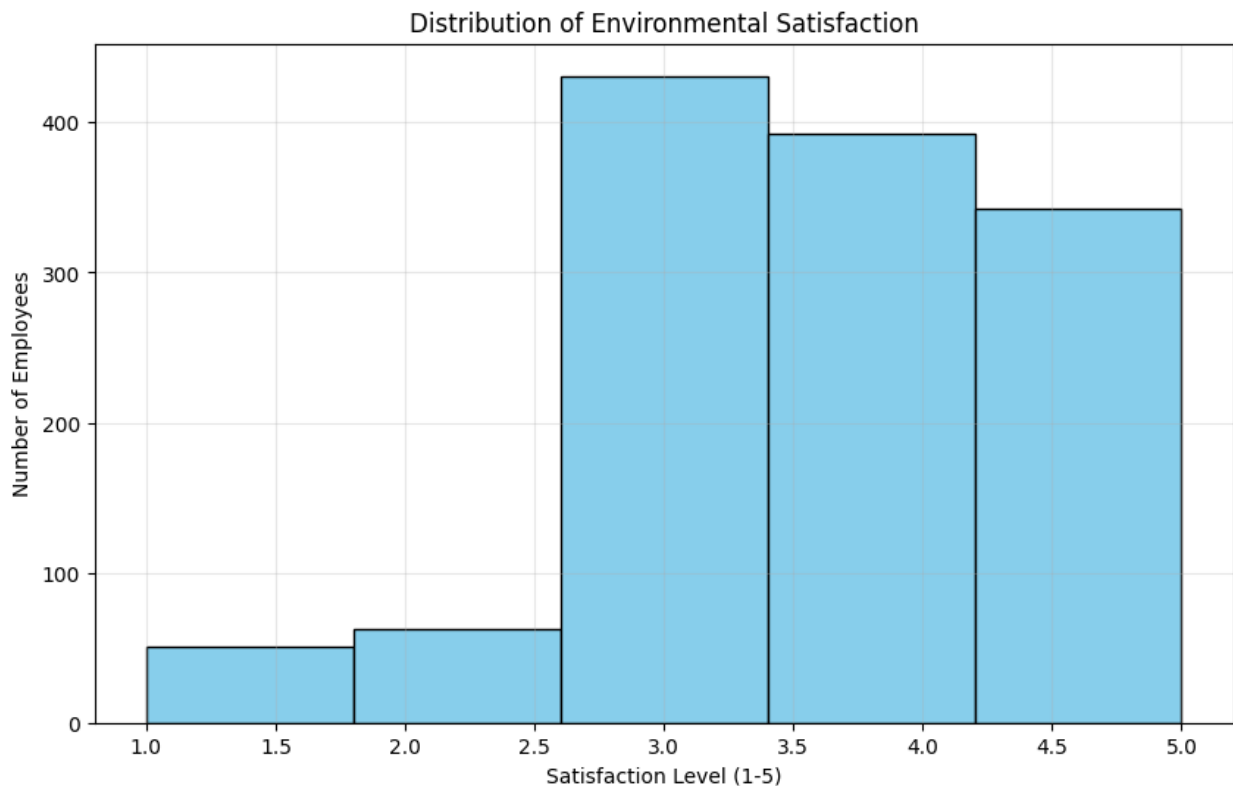
plt.ylabel('Number of Employees')
plt.grid(True, alpha=0.3)
plt.show()

# Calculate correlation with ManagerRating
correlation =
merged_df['EnvironmentSatisfaction'].corr(merged_df['ManagerRating'])
print(f"Correlation between Environmental Satisfaction and Manager
Rating: {correlation:.2f}")

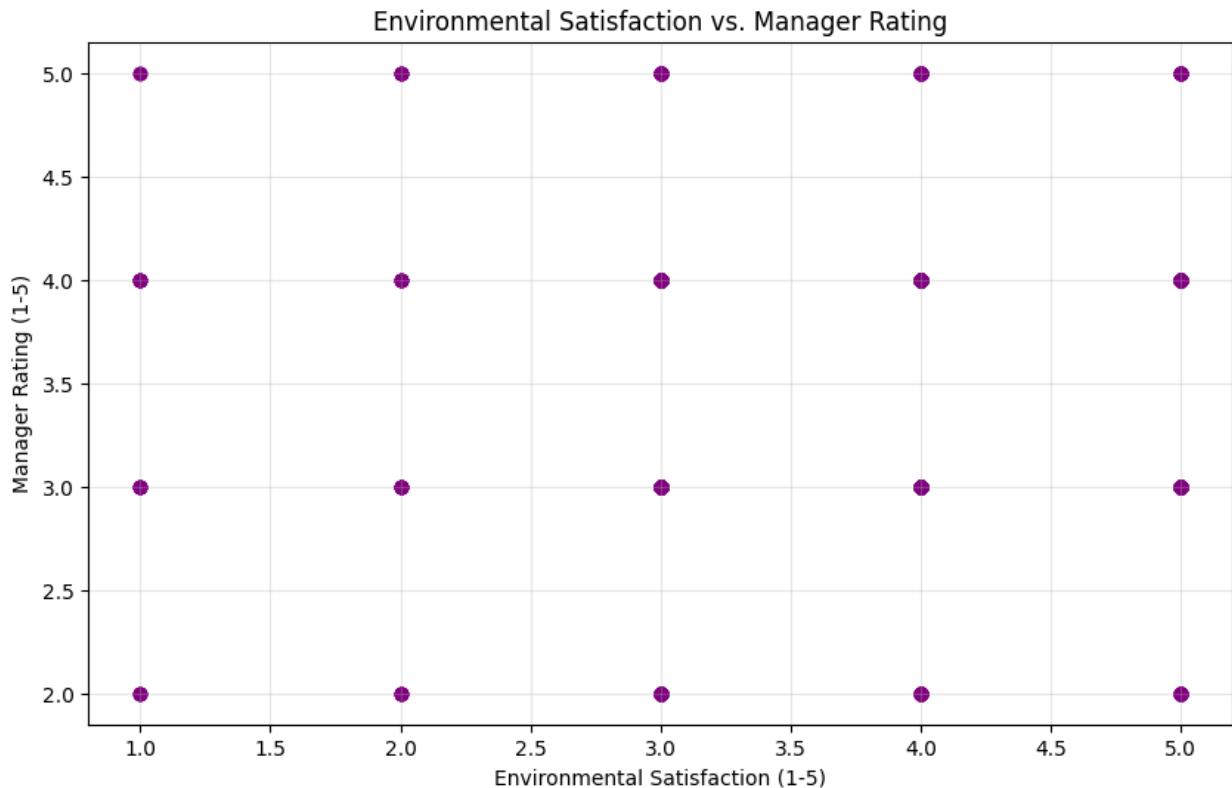
# Scatter plot of Environmental Satisfaction vs. ManagerRating
plt.figure(figsize=(10, 6))
plt.scatter(merged_df['EnvironmentSatisfaction'],
merged_df['ManagerRating'], alpha=0.5, color='purple')
plt.title('Environmental Satisfaction vs. Manager Rating')
plt.xlabel('Environmental Satisfaction (1-5)')
plt.ylabel('Manager Rating (1-5)')
plt.grid(True, alpha=0.3)
plt.show()

```

Average Environmental Satisfaction (1-5 scale): 3.71



Correlation between Environmental Satisfaction and Manager Rating:
0.02



Interpretation:

- **Average Environmental Satisfaction:** A score near 5 indicates high satisfaction, while a score below 3 suggests dissatisfaction and a need for improvement in the work environment or culture.
 - **Distribution:** If the histogram skews left (higher values), most employees are satisfied. A right skew (lower values) indicates widespread dissatisfaction.
 - **Correlation:** A positive value (e.g., 0.3 or higher) suggests that satisfied employees tend to receive better performance ratings, reinforcing the link between engagement and performance.
-

Analysis 2: Predictive Analytics

Objective: This analysis uses machine learning to predict human resource trends:

1. Which employees are most likely to leave? (Attrition prediction)
2. What factors influence high performance? (Performance prediction)
3. How can recruitment strategies be improved? (Inferred from insights)

1. Attrition Prediction

```
# Prepare data for attrition prediction
attrition_df = merged_df.copy()
attrition_df['Attrition'] = attrition_df['Attrition'].map({'Yes': 1,
'No': 0})
```

```

# Select features for prediction
features = ['Age', 'DistanceFromHome', 'EnvironmentSatisfaction',
            'JobSatisfaction',
            'WorkLifeBalance', 'YearsAtCompany',
            'YearsSinceLastPromotion', 'ManagerRating']

# Remove rows with missing values
attrition_df = attrition_df.dropna(subset=features + ['Attrition'])

# Split data into training and testing sets
X = attrition_df[features]
y = attrition_df['Attrition']
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.3, random_state=42)

# Train a Random Forest Classifier
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)

# Predict and evaluate
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Attrition Prediction Accuracy: {accuracy:.2f}")
print("Classification Report:\n", classification_report(y_test,
y_pred))

# Feature importance
importances = clf.feature_importances_
feature_importance_df = pd.DataFrame({'Feature': features,
'Importance': importances})
feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)
print("Feature Importances for Attrition Prediction:\n",
feature_importance_df)

# Visualize feature importances
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df,
palette='Reds_d')
plt.title('Feature Importances for Attrition Prediction')
plt.show()

```

Attrition Prediction Accuracy: 0.89

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.98	0.94	314
1	0.85	0.50	0.63	70
accuracy			0.89	384

macro avg	0.88	0.74	0.78	384
weighted avg	0.89	0.89	0.88	384

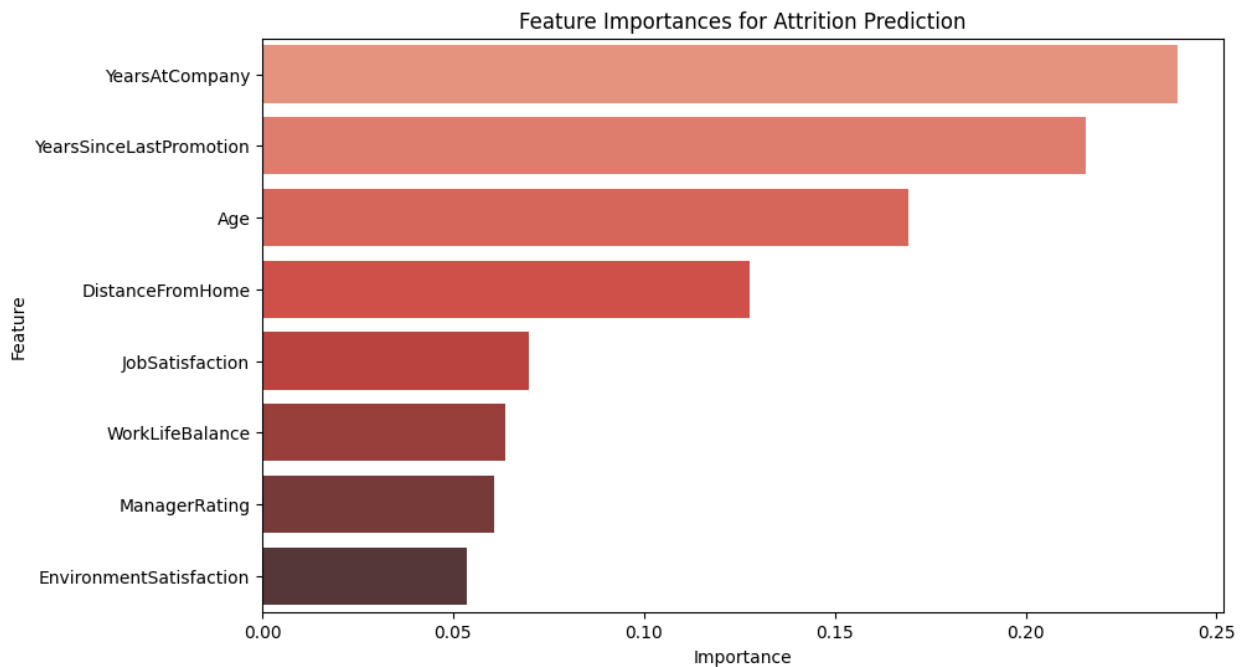
Feature Importances for Attrition Prediction:

	Feature	Importance
5	YearsAtCompany	0.239835
6	YearsSinceLastPromotion	0.215694
0	Age	0.169258
1	DistanceFromHome	0.127701
3	JobSatisfaction	0.069722
4	WorkLifeBalance	0.063582
7	ManagerRating	0.060758
2	EnvironmentSatisfaction	0.053449

/tmp/ipykernel_160157/684693721.py:35: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=feature_importance_df,
palette='Reds_d')
```



2. Performance Prediction

```
# Prepare data for performance prediction
performance_df = merged_df.copy()
```

```
# Select features for prediction
```

```

features = ['Age', 'EnvironmentSatisfaction', 'JobSatisfaction',
            'WorkLifeBalance',
            'YearsAtCompany', 'ManagerRating']

# Remove rows with missing values
performance_df = performance_df.dropna(subset=features)

# Split data into training and testing sets
X = performance_df[features]
y = performance_df['ManagerRating']
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.3, random_state=42)

# Train a Random Forest Regressor
reg = RandomForestRegressor(random_state=42)
reg.fit(X_train, y_train)

# Predict and evaluate
y_pred = reg.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Performance Prediction Mean Squared Error: {mse:.2f}")

# Feature importance
importances = reg.feature_importances_
feature_importance_df = pd.DataFrame({'Feature': features,
                                     'Importance': importances})
feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)
print("Feature Importances for Performance Prediction:\n",
      feature_importance_df)

# Visualize feature importances
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df,
           palette='Purples_d')
plt.title('Feature Importances for Performance Prediction')
plt.show()

```

Performance Prediction Mean Squared Error: 0.00

Feature Importances for Performance Prediction:

	Feature	Importance
5	ManagerRating	1.0
0	Age	0.0
1	EnvironmentSatisfaction	0.0
2	JobSatisfaction	0.0
3	WorkLifeBalance	0.0
4	YearsAtCompany	0.0

/tmp/ipykernel_160157/3741518954.py:33: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=feature_importance_df, palette='Purples_d')
```

